

DATA245_ForestFireSizePrediction_GRP7

November 13, 2023

1 Forest Fire Size Prediction

1.0.1 Importing all the required libraries.

```
[1]: import warnings
      warnings.filterwarnings('ignore')

[2]: import numpy as np
      import matplotlib.pyplot as plt
      import pandas as pd
      import seaborn as sns
      from datetime import datetime
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import OneHotEncoder
      from sklearn.linear_model import LogisticRegression
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.svm import SVC
      from sklearn.naive_bayes import GaussianNB
      import category_encoders as ce
```

1.0.2 Reading the CSV files

```
[3]: df = pd.read_csv("data/FW_Veg_Rem_Combined.csv")
      df.head(5)
```

```
[3]:   Unnamed: 0.1  Unnamed: 0  fire_name  fire_size  fire_size_class  \
0              0            0      NaN         10.0                C
1              1            1      NaN          3.0                B
2              2            2      NaN         60.0                C
3              3            3    WNA  1          1.0                B
4              4            4      NaN          2.0                B

      stat_cause_descr  latitude  longitude  state  disc_clean_date  ...  \
0  Missing/Undefined  18.105072  -66.753044    PR         2/11/2007  ...
1              Arson  35.038330  -87.610000    TN         12/11/2006  ...
```

2	Arson	34.947800	-88.722500	MS	2/29/2004	...
3	Debris Burning	39.641400	-119.308300	NV	6/6/2005	...
4	Miscellaneous	30.700600	-90.591400	LA	9/22/1999	...

	Wind_cont	Hum_pre_30	Hum_pre_15	Hum_pre_7	Hum_cont	Prec_pre_30	\
0	3.250413	78.216590	76.793750	76.381579	78.724370	0.0	
1	2.122320	70.840000	65.858911	55.505882	81.682678	59.8	
2	3.369050	75.531629	75.868613	76.812834	65.063800	168.8	
3	0.000000	44.778429	37.140811	35.353846	0.000000	10.4	
4	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.0	

	Prec_pre_15	Prec_pre_7	Prec_cont	remoteness
0	0.0	0.0	0.0	0.017923
1	8.4	0.0	86.8	0.184355
2	42.2	18.1	124.5	0.194544
3	7.2	0.0	0.0	0.487447
4	-1.0	-1.0	-1.0	0.214633

[5 rows x 43 columns]

```
[4]: df.describe(include = "all")
```

```
[4]:
```

	Unnamed: 0.1	Unnamed: 0	fire_name	fire_size	fire_size_class	\
count	55367.000000	55367.000000	25913	55367.000000	55367	
unique	NaN	NaN	21793	NaN	6	
top	NaN	NaN	GRASS FIRE	NaN	B	
freq	NaN	NaN	128	NaN	36522	
mean	27683.000000	27683.000000	NaN	2104.645161	NaN	
std	15983.220514	15983.220514	NaN	14777.005364	NaN	
min	0.000000	0.000000	NaN	0.510000	NaN	
25%	13841.500000	13841.500000	NaN	1.200000	NaN	
50%	27683.000000	27683.000000	NaN	4.000000	NaN	
75%	41524.500000	41524.500000	NaN	20.000000	NaN	
max	55366.000000	55366.000000	NaN	606945.000000	NaN	

	stat_cause_descr	latitude	longitude	state	disc_clean_date	\
count	55367	55367.000000	55367.000000	55367	55367	
unique	13	NaN	NaN	51	8114	
top	Debris Burning	NaN	NaN	TX	6/21/2008	
freq	14278	NaN	NaN	6080	64	
mean	NaN	36.172866	-94.757971	NaN	NaN	
std	NaN	6.724348	15.878194	NaN	NaN	
min	NaN	17.956533	-165.936000	NaN	NaN	
25%	NaN	32.265960	-102.541513	NaN	NaN	
50%	NaN	34.600000	-91.212359	NaN	NaN	
75%	NaN	38.975235	-82.847500	NaN	NaN	
max	NaN	69.849500	-65.285833	NaN	NaN	

	...	Wind_cont	Hum_pre_30	Hum_pre_15	Hum_pre_7	\
count	...	55367.000000	55367.000000	55367.000000	55367.000000	
unique	...	NaN	NaN	NaN	NaN	
top	...	NaN	NaN	NaN	NaN	
freq	...	NaN	NaN	NaN	NaN	
mean	...	1.132284	40.781796	38.453935	37.001865	
std	...	2.030611	31.086856	31.042541	30.827885	
min	...	-1.000000	-1.000000	-1.000000	-1.000000	
25%	...	-1.000000	-1.000000	-1.000000	-1.000000	
50%	...	0.000000	55.657480	51.753846	48.230769	
75%	...	2.848603	67.384352	65.911469	64.645296	
max	...	24.200000	96.000000	94.000000	96.000000	

		Hum_cont	Prec_pre_30	Prec_pre_15	Prec_pre_7	Prec_cont	\
count	55367.000000	55367.000000	55367.000000	55367.000000	55367.000000	55367.000000	
unique	NaN	NaN	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	NaN	NaN	
mean	25.056738	26.277046	11.654253	4.689920	15.590440		
std	31.187638	112.050198	56.920510	31.205327	59.757113		
min	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000		
25%	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000		
50%	0.000000	0.000000	0.000000	0.000000	0.000000		
75%	60.193606	18.900000	3.600000	0.000000	0.000000		
max	94.000000	13560.800000	2527.000000	1638.000000	2126.000000		

	remoteness
count	55367.000000
unique	NaN
top	NaN
freq	NaN
mean	0.236799
std	0.144865
min	0.000000
25%	0.137800
50%	0.202114
75%	0.284782
max	1.000000

[11 rows x 43 columns]

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55367 entries, 0 to 55366
Data columns (total 43 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0.1	55367 non-null	int64
1	Unnamed: 0	55367 non-null	int64
2	fire_name	25913 non-null	object
3	fire_size	55367 non-null	float64
4	fire_size_class	55367 non-null	object
5	stat_cause_descr	55367 non-null	object
6	latitude	55367 non-null	float64
7	longitude	55367 non-null	float64
8	state	55367 non-null	object
9	disc_clean_date	55367 non-null	object
10	cont_clean_date	27477 non-null	object
11	discovery_month	55367 non-null	object
12	disc_date_final	28708 non-null	object
13	cont_date_final	25632 non-null	object
14	putout_time	27477 non-null	object
15	disc_date_pre	55367 non-null	object
16	disc_pre_year	55367 non-null	int64
17	disc_pre_month	55367 non-null	object
18	wstation_usaf	55367 non-null	object
19	dstation_m	55367 non-null	float64
20	wstation_wban	55367 non-null	int64
21	wstation_byear	55367 non-null	int64
22	wstation_eyear	55367 non-null	int64
23	Vegetation	55367 non-null	int64
24	fire_mag	55367 non-null	float64
25	weather_file	55367 non-null	object
26	Temp_pre_30	55367 non-null	float64
27	Temp_pre_15	55367 non-null	float64
28	Temp_pre_7	55367 non-null	float64
29	Temp_cont	55367 non-null	float64
30	Wind_pre_30	55367 non-null	float64
31	Wind_pre_15	55367 non-null	float64
32	Wind_pre_7	55367 non-null	float64
33	Wind_cont	55367 non-null	float64
34	Hum_pre_30	55367 non-null	float64
35	Hum_pre_15	55367 non-null	float64
36	Hum_pre_7	55367 non-null	float64
37	Hum_cont	55367 non-null	float64
38	Prec_pre_30	55367 non-null	float64
39	Prec_pre_15	55367 non-null	float64
40	Prec_pre_7	55367 non-null	float64
41	Prec_cont	55367 non-null	float64
42	remoteness	55367 non-null	float64

dtypes: float64(22), int64(7), object(14)

memory usage: 18.2+ MB

Data Dictionary

- Fire_size_class - Class of Fire Size (A-G)
- Stat_cause_descr - Cause of Fire
- Latitude - Latitude of Fire
- Longitude - Longitude of Fire
- Discovery_month - Month in which Fire was discovered
- Vegetation - Dominant vegetation in the areas (can save some factors of vegetation)
- Temp_pre - temperature in deg C at the location of fire up to 30, 15 and 7 days prior
- Temp_cont - temperature in deg C at the location of fire up to day the fire was
- Wind_pre - wind in deg C at the location of fire up to 30, 15 and 7 days prior
- Wind_cont - wind in deg C at the location of fire up to day the fire was
- Prec_pre - Precipitation in deg C at the location of fire up to 30, 15 and 7 days prior
- Prec_cont - Precipitation in deg C at the location of fire up to day the fire was
- Hum_pre - Humidity in deg C at the location of fire up to 30, 15 and 7 days prior
- Hum_cont - Humidity in deg C at the location of fire up to day the fire was
- Remoteness - non-dimensional distance to closest city

1.0.3 Removing the redundant and unnecessary columns

```
[6]: df = df.drop(['Unnamed: 0.1', 'Unnamed: 0', 'fire_name', 'state',  
↳ 'cont_clean_date',  
↳ 'discovery_month', 'disc_date_final', 'cont_date_final',  
↳ 'putout_time', 'disc_pre_year', 'disc_pre_month',  
↳ 'wstation_usaf', 'dstation_m', 'wstation_wban', 'fire_mag',  
↳ 'weather_file'],axis=1)
```

Removed columns with null values, redundant columns like fire_mag, fire_size and date variables

```
[7]: df.head(5)
```

```
[7]:
```

	fire_size	fire_size_class	stat_cause_descr	latitude	longitude	\
0	10.0	C	Missing/Undefined	18.105072	-66.753044	
1	3.0	B	Arson	35.038330	-87.610000	
2	60.0	C	Arson	34.947800	-88.722500	
3	1.0	B	Debris Burning	39.641400	-119.308300	
4	2.0	B	Miscellaneous	30.700600	-90.591400	

	disc_clean_date	disc_date_pre	wstation_byear	wstation_eyear	Vegetation	\
0	2/11/2007	1/12/2007	1945	2018	12	
1	12/11/2006	11/11/2006	1978	2020	15	
2	2/29/2004	1/30/2004	1978	2020	16	
3	6/6/2005	5/7/2005	1942	2020	0	
4	9/22/1999	8/23/1999	1987	2016	12	

	...	Wind_cont	Hum_pre_30	Hum_pre_15	Hum_pre_7	Hum_cont	Prec_pre_30	\
0	...	3.250413	78.216590	76.793750	76.381579	78.724370	0.0	

1	...	2.122320	70.840000	65.858911	55.505882	81.682678	59.8
2	...	3.369050	75.531629	75.868613	76.812834	65.063800	168.8
3	...	0.000000	44.778429	37.140811	35.353846	0.000000	10.4
4	...	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.0

	Prec_pre_15	Prec_pre_7	Prec_cont	remoteness
0	0.0	0.0	0.0	0.017923
1	8.4	0.0	86.8	0.184355
2	42.2	18.1	124.5	0.194544
3	7.2	0.0	0.0	0.487447
4	-1.0	-1.0	-1.0	0.214633

[5 rows x 27 columns]

1.0.4 Data Visualization

```
[8]: import matplotlib.pyplot as plt
import seaborn as sns
import folium
from folium.plugins import MarkerCluster, HeatMap
from folium import Choropleth
```

```
[9]: # Create a base map centered at a specific location
map_center = [df['latitude'].mean(), df['longitude'].mean()]
mymap = folium.Map(location=map_center, zoom_start=4)

# Create a MarkerCluster to group markers at the same location
marker_cluster = MarkerCluster().add_to(mymap)

# Add markers for each data point
for index, row in df.iterrows():
    folium.Marker(
        location=[row['latitude'], row['longitude']],
        popup=f"Fire Size Class: {row['fire_size_class']}",
        icon=None, # You can customize the icon if needed
    ).add_to(marker_cluster)

# Display the map in the notebook
display(mymap)
```

<folium.folium.Map at 0x7fbeb9183220>

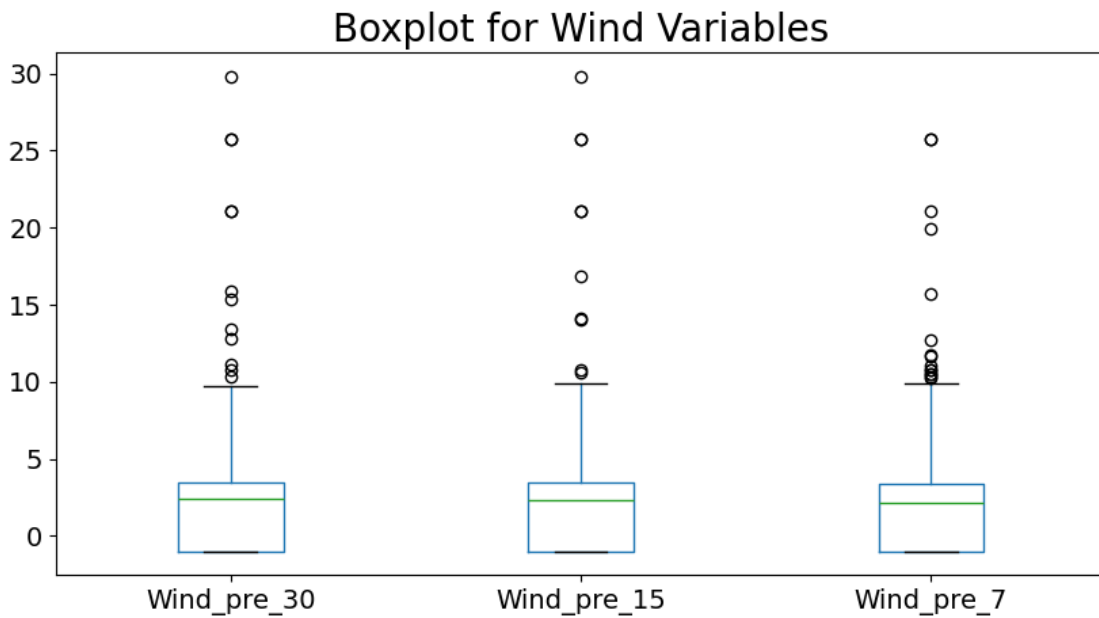
The dataset has fire occurrences data from all over US, concentrated in California and Southeastern part.

```
[10]: plt.figure(figsize=(10, 5))
plt.title("Boxplot for Wind Variables", fontsize=20)
plt.grid(color='grey', linestyle='-', linewidth=0.25, alpha=0.5)
```

```
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)

# Create boxplot
df.boxplot(column=["Wind_pre_30", "Wind_pre_15", "Wind_pre_7"], grid=False)

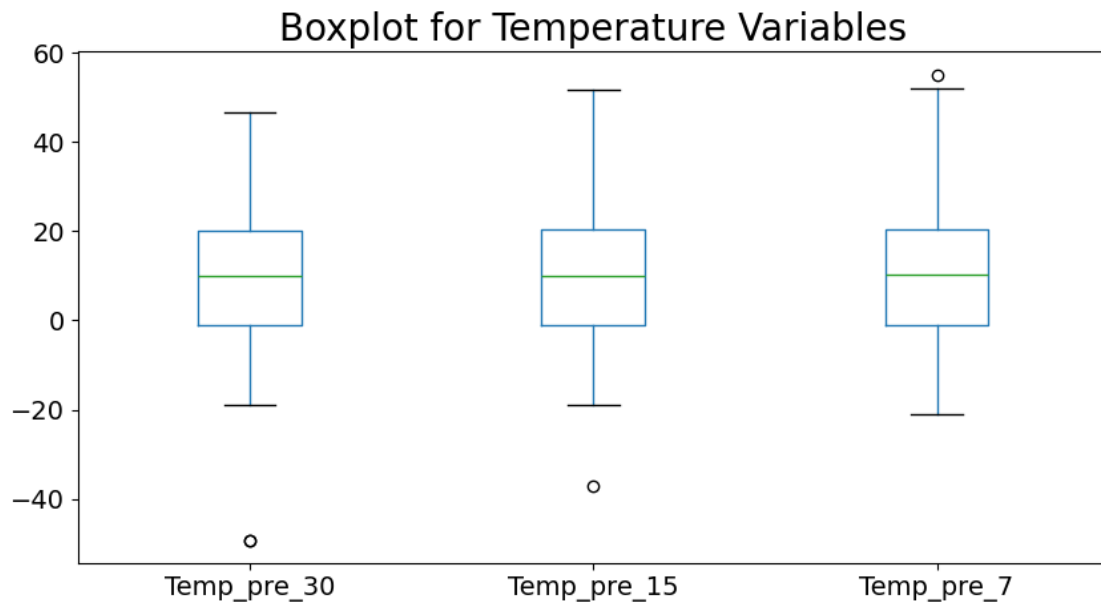
plt.show()
```



```
[11]: plt.figure(figsize=(10, 5))
plt.title("Boxplot for Temperature Variables", fontsize=20)
plt.grid(color='grey', linestyle='--', linewidth=0.25, alpha=0.5)
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)

# Create boxplot
df.boxplot(column=["Temp_pre_30", "Temp_pre_15", "Temp_pre_7"], grid=False)


plt.show()
```



```
[12]: plt.figure(figsize=(10, 5))
plt.title("Boxplot for Humidity Variables", fontsize=20)
plt.grid(color='grey', linestyle='--', linewidth=0.25, alpha=0.5)
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)

# Create boxplot
df.boxplot(column=["Hum_pre_30", "Hum_pre_15", "Hum_pre_7"], grid=False)

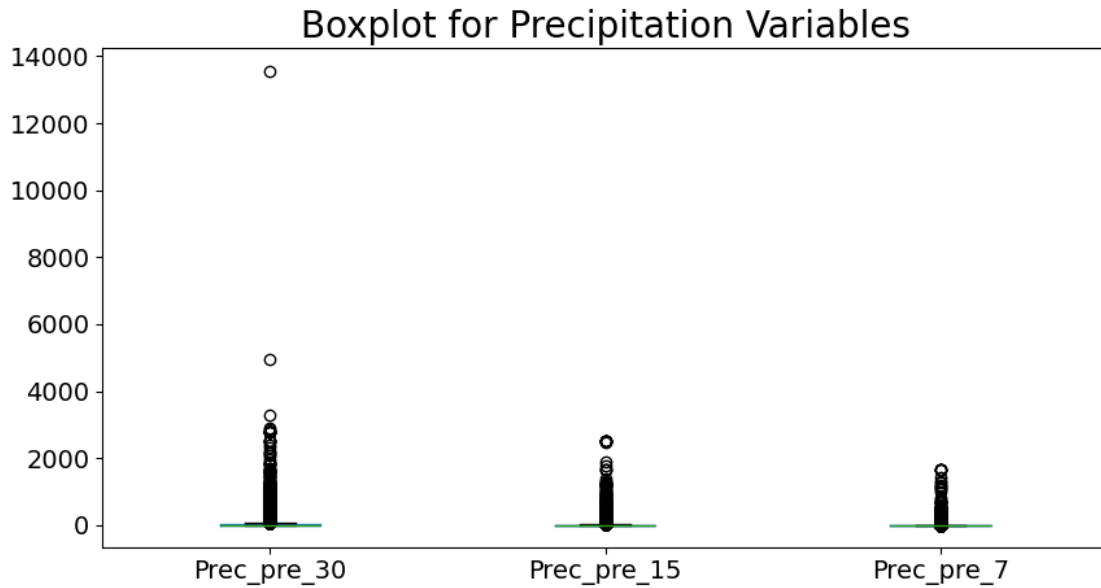
plt.show()
```

```
[13]: plt.figure(figsize=(10, 5))
plt.title("Boxplot for Precipitation Variables", fontsize=20)
plt.grid(color='grey', linestyle='--', linewidth=0.25, alpha=0.5)
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)

# Create boxplot
df.boxplot(column=["Prec_pre_30", "Prec_pre_15", "Prec_pre_7"], grid=False)

plt.show()
```



Will apply minmaxscaler to take care of these outliers in the weather data

```
[14]: (df.fire_size_class.value_counts()/df.shape[0])*100
```

```
[14]: fire_size_class
B      65.963480
C      19.526071
G       7.173948
F       3.554464
D       2.517745
E       1.264291
Name: count, dtype: float64
```

As the target class is unbalanced we decided to club the smaller group of classes as 1.

Clubbing (C,D,E,F,G) as (1) class and A,B as (0) class.

0 indicates small fire <25 Acres and 1 represents a widespread fire >25Acres.

```
[15]: class_mapping = {'A': 0, 'B': 0, 'C':1, 'D':1, 'E':1, 'F':1, 'G':1}
df = df.replace(class_mapping)
```

```
[16]: df.fire_size_class.value_counts()
```

```
[16]: fire_size_class
0      36522
1      18845
Name: count, dtype: int64
```

```
[17]: (df.fire_size_class.value_counts()/df.shape[0])*100
```

```
[17]: fire_size_class
0    65.96348
1    34.03652
Name: count, dtype: float64
```

1.0.5 Extracting the date and month and removing the redundant columns.


```
[18]: # Extract day, month, year from discovery clean date
df['disc_clean_date'] = pd.to_datetime(df['disc_clean_date'])

df['disc_month'] = df['disc_clean_date'].dt.month

# Drop the columns which are not required
df = df.drop(['disc_clean_date', 'disc_date_pre', \
             'wstation_byear', 'wstation_eyear'],axis=1)
```

```
[19]: # Group by 'disc_month' and calculate the average fire size for each month
average_fire_size_per_month = df.groupby('disc_month')['fire_size'].mean().
    ↪reset_index()

# Create a bar chart
plt.figure(figsize=(10, 6))
plt.bar(average_fire_size_per_month['disc_month'],
    ↪average_fire_size_per_month['fire_size'], color='red')
plt.xlabel('Month')
plt.ylabel('Average Fire Size')
plt.title('Average Fire Size by Month')
plt.show()
```



1.0.6 Bigger fires occur during the summer season, so months columns will be an important variable

```
[20]: df= df.drop(["fire_size"], axis = 1)
```

```
[21]: df['Vegetation'] = df['Vegetation'].astype(object)
```

1.0.7 Applying MinMaxScaler to the weather variables

```
[22]: from sklearn.preprocessing import MinMaxScaler

trans = MinMaxScaler()
df.iloc[:, 5:21] = trans.fit_transform(df.iloc[:, 5:21])
```

```
[23]: X = df.drop('fire_size_class',axis=1)
y = df['fire_size_class']
```

```
[24]: X.columns
```

```
[24]: Index(['stat_cause_descr', 'latitude', 'longitude', 'Vegetation',
          'Temp_pre_30', 'Temp_pre_15', 'Temp_pre_7', 'Temp_cont', 'Wind_pre_30',
          'Wind_pre_15', 'Wind_pre_7', 'Wind_cont', 'Hum_pre_30', 'Hum_pre_15',
          'Hum_pre_7', 'Hum_cont', 'Prec_pre_30', 'Prec_pre_15', 'Prec_pre_7',
```

```

    'Prec_cont', 'remoteness', 'disc_month'],
    dtype='object')

```

1.0.8 Defining the function for target encoding

```

[25]: def target_encode_multiclass(X,y): #X,y are pandas df and series
      y=y.astype(str) #convert to string to onehot encode
      enc=ce.OneHotEncoder().fit(y)
      y_onehot=enc.transform(y)
      class_names=y_onehot.columns #names of onehot encoded columns
      X_obj=X.select_dtypes('object') #separate categorical columns
      X=X.select_dtypes(exclude='object')
      for class_ in class_names:
          enc=ce.TargetEncoder()
          enc.fit(X_obj,y_onehot[class_]) #convert all categorical
          temp=enc.transform(X_obj) #columns for class_
          temp.columns=[str(x)+'_'+str(class_) for x in temp.columns]
          X=pd.concat([X,temp],axis=1) #add to original dataset

      return X

```

```

[26]: X = target_encode_multiclass(X,y)

```

```

[27]: X.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55367 entries, 0 to 55366
Data columns (total 24 columns):

```


#	Column	Non-Null Count	Dtype
0	latitude	55367 non-null	float64
1	longitude	55367 non-null	float64
2	Temp_pre_30	55367 non-null	float64
3	Temp_pre_15	55367 non-null	float64
4	Temp_pre_7	55367 non-null	float64
5	Temp_cont	55367 non-null	float64
6	Wind_pre_30	55367 non-null	float64
7	Wind_pre_15	55367 non-null	float64
8	Wind_pre_7	55367 non-null	float64
9	Wind_cont	55367 non-null	float64
10	Hum_pre_30	55367 non-null	float64
11	Hum_pre_15	55367 non-null	float64
12	Hum_pre_7	55367 non-null	float64
13	Hum_cont	55367 non-null	float64
14	Prec_pre_30	55367 non-null	float64
15	Prec_pre_15	55367 non-null	float64
16	Prec_pre_7	55367 non-null	float64

```

17 Prec_cont                55367 non-null float64
18 remoteness              55367 non-null float64
19 disc_month              55367 non-null int32
20 stat_cause_descr_fire_size_class_1 55367 non-null float64
21 Vegetation_fire_size_class_1      55367 non-null float64
22 stat_cause_descr_fire_size_class_2 55367 non-null float64
23 Vegetation_fire_size_class_2      55367 non-null float64
dtypes: float64(23), int32(1)
memory usage: 9.9 MB

```

```
[28]: sns.heatmap(pd.concat([X, y], axis = 1).corr())
plt.show()
```



```
[29]: X.describe(include="all")
```

```

[29]:
count    latitude    longitude    Temp_pre_30    Temp_pre_15    Temp_pre_7  \
mean      36.172866    -94.757971      0.625182      0.539708      0.418654

```

std	6.724348	15.878194	0.109253	0.120064	0.142499
min	17.956533	-165.936000	0.000000	0.000000	0.000000
25%	32.265960	-102.541513	0.503186	0.406468	0.261663
50%	34.600000	-91.212359	0.617428	0.532045	0.409829
75%	38.975235	-82.847500	0.722531	0.646503	0.545169
max	69.849500	-65.285833	1.000000	1.000000	1.000000

	Temp_cont	Wind_pre_30	Wind_pre_15	Wind_pre_7	Wind_cont \
count	55367.000000	55367.000000	55367.000000	55367.000000	55367.000000
mean	0.393417	0.095004	0.092342	0.104544	0.084614
std	0.135891	0.068382	0.068919	0.080139	0.080580
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.279743	0.000000	0.000000	0.000000	0.000000
50%	0.292605	0.111132	0.107011	0.118930	0.039683
75%	0.511406	0.145122	0.144009	0.164504	0.152722
max	1.000000	1.000000	1.000000	1.000000	1.000000

	...	Prec_pre_30	Prec_pre_15	Prec_pre_7	Prec_cont \
count	...	55367.000000	55367.000000	55367.000000	55367.000000
mean	...	0.002011	0.005006	0.003472	0.007800
std	...	0.008262	0.022516	0.019039	0.028095
min	...	0.000000	0.000000	0.000000	0.000000
25%	...	0.000000	0.000000	0.000000	0.000000
50%	...	0.000074	0.000396	0.000610	0.000470
75%	...	0.001467	0.001820	0.000610	0.000470
max	...	1.000000	1.000000	1.000000	1.000000

	remoteness	disc_month	stat_cause_descr_fire_size_class_1 \
count	55367.000000	55367.000000	55367.000000
mean	0.236799	5.694331	0.340366
std	0.144865	3.024138	0.136714
min	0.000000	1.000000	0.133594
25%	0.137800	3.000000	0.224751
50%	0.202114	5.000000	0.329925
75%	0.284782	8.000000	0.349959
max	1.000000	12.000000	0.639328

	Vegetation_fire_size_class_1	stat_cause_descr_fire_size_class_2 \
count	55367.000000	55367.000000
mean	0.340365	0.659634
std	0.052303	0.136714
min	0.298379	0.360672
25%	0.298379	0.650041
50%	0.299193	0.670075
75%	0.399263	0.775249
max	0.440397	0.866406

	Vegetation_fire_size_class_2
count	55367.000000
mean	0.659635
std	0.052303
min	0.559603
25%	0.600737
50%	0.700807
75%	0.701621
max	0.701621

[8 rows x 24 columns]

```
[30]: from sklearn.metrics import roc_curve, accuracy_score
from sklearn.metrics import auc, classification_report
from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedKFold
```

1.0.9 Calculating Accuracy for 5 folds

```
[31]: X = X.values
y = y.values

# Define the number of folds for cross-validation
num_folds = 5

# Initialize StratifiedKFold for stratified sampling based on the distribution
# of classes
stratified_kfold = StratifiedKFold(n_splits=num_folds,
                                   shuffle=True,
                                   random_state=42)
```

```
[32]: classifiers=[('Logistic Regression :', LogisticRegression()),
                  ('Decision Tree Classification :', DecisionTreeClassifier()),
                  ('Random Forest Classification :', RandomForestClassifier()),
                  ('K-Neighbors Classification :', KNeighborsClassifier()),
                  ('Gaussian Naive Bayes :', GaussianNB()),
                  ('Support Vector Classification :', SVC())]

cla_pred=[]

for name,model in classifiers:
    model=model
    # Perform K-fold cross-validation
    for fold, (train_index, test_index) in enumerate(stratified_kfold.split(X,
    y)):
        X_train, X_test = X[train_index], X[test_index]
```



```

y_train, y_test = y[train_index], y[test_index]

model.fit(X_train, y_train)

# Make predictions on the test set
predictions = model.predict(X_test)

# Evaluate the model
cla_pred.append(accuracy_score(y_test,predictions))
print(name, fold+1, accuracy_score(y_test,predictions))

```

```

Logistic Regression : 1 0.7214195412678346
Logistic Regression : 2 0.7254831135994221
Logistic Regression : 3 0.725729251332069
Logistic Regression : 4 0.7242842951323039
Logistic Regression : 5 0.7306962882687619
Decision Tree Classification : 1 0.6708506411414124
Decision Tree Classification : 2 0.679971103485642
Decision Tree Classification : 3 0.6780456967398176
Decision Tree Classification : 4 0.6751557843402872
Decision Tree Classification : 5 0.6796712724645534
Random Forest Classification : 1 0.7515802781289507
Random Forest Classification : 2 0.7519414845584251
Random Forest Classification : 3 0.7520093922152985
Random Forest Classification : 4 0.7546283753273729
Random Forest Classification : 5 0.7567958096270206
K-Neighbors Classification : 1 0.7165432544699296
K-Neighbors Classification : 2 0.7174462705436156
K-Neighbors Classification : 3 0.7101959721845932
K-Neighbors Classification : 4 0.715704867696198
K-Neighbors Classification : 5 0.7117312381468437
Gaussian Naive Bayes : 1 0.6743724038287882
Gaussian Naive Bayes : 2 0.6794292938414304
Gaussian Naive Bayes : 3 0.6756073331527138
Gaussian Naive Bayes : 4 0.6798518919895241
Gaussian Naive Bayes : 5 0.677052289352479
Support Vector Classification : 1 0.6731984829329962
Support Vector Classification : 2 0.6744627054361567
Support Vector Classification : 3 0.673981757427978
Support Vector Classification : 4 0.673710828140522
Support Vector Classification : 5 0.6749751648153165

```

1.0.10 Printing Classification Report for 5 folds

Logistic Regression

```

[33]: X1 = df.drop('fire_size_class',axis=1)
      y1 = df['fire_size_class']

```

```

X1 = target_encode_multiclass(X1, y1)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X1, y1, test_size=0.2,
    ↪random_state=42)

# Set up k-fold cross-validation
stratified_kf = StratifiedKFold(n_splits=num_folds, shuffle=True,
    ↪random_state=42)

```

```

[34]: LR = LogisticRegression()

# Lists to store metrics for each fold
LR_precision_list = []
LR_recall_list = []
LR_f1_list = []

for train_index, val_index in stratified_kf.split(X_train, y_train):
    X_train_fold, X_val_fold = X_train.iloc[train_index], X_train.
    ↪iloc[val_index]
    y_train_fold, y_val_fold = y_train.iloc[train_index], y_train.
    ↪iloc[val_index]

    LR.fit(X_train_fold, y_train_fold)
    LR_pred = LR.predict(X_val_fold)

    # Calculate metrics for each fold
    classification_report_fold = classification_report(y_val_fold, LR_pred,
    ↪output_dict=True)

    LR_precision_list.append(classification_report_fold['weighted_
    ↪avg']['precision'])
    LR_recall_list.append(classification_report_fold['weighted avg']['recall'])
    LR_f1_list.append(classification_report_fold['weighted avg']['f1-score'])

# Calculate mean metrics across all folds
mean_LR_precision = np.mean(LR_precision_list)
mean_LR_recall = np.mean(LR_recall_list)
mean_LR_f1 = np.mean(LR_f1_list)

# Print or use the mean metrics as needed
print(f'Mean Precision for LR: {mean_LR_precision}')
print(f'Mean Recall for LR: {mean_LR_recall}')
print(f'Mean F1-Score for LR: {mean_LR_f1}')

```

Mean Precision for LR: 0.7175383891667414

Mean Recall for LR: 0.7224844405762785
Mean F1-Score for LR: 0.6879843015053047

DecisionTree

```
[35]: DT = DecisionTreeClassifier()

# Lists to store metrics for each fold
DT_precision_list = []
DT_recall_list = []
DT_f1_list = []

for train_index, val_index in stratified_kf.split(X_train, y_train):
    X_train_fold, X_val_fold = X_train.iloc[train_index], X_train.
    ↪iloc[val_index]
    y_train_fold, y_val_fold = y_train.iloc[train_index], y_train.
    ↪iloc[val_index]

    DT.fit(X_train_fold, y_train_fold)
    DT_pred = DT.predict(X_val_fold)

    # Calculate metrics for each fold
    classification_report_fold = classification_report(y_val_fold, DT_pred,
    ↪output_dict=True)

    DT_precision_list.append(classification_report_fold['weighted_
    ↪avg']['precision'])
    DT_recall_list.append(classification_report_fold['weighted avg']['recall'])
    DT_f1_list.append(classification_report_fold['weighted avg']['f1-score'])

# Calculate mean metrics across all folds
mean_DT_precision = np.mean(DT_precision_list)
mean_DT_recall = np.mean(DT_recall_list)
mean_DT_f1 = np.mean(DT_f1_list)

# Print or use the mean metrics as needed
print(f'Mean Precision for DT: {mean_DT_precision}')
print(f'Mean Recall for DT: {mean_DT_recall}')
print(f'Mean F1-Score for DT: {mean_DT_f1}')
```

Mean Precision for DT: 0.6809505451590552
Mean Recall for DT: 0.6792944892577222
Mean F1-Score for DT: 0.6800503067789476

Random Forest

```
[36]: RF = RandomForestClassifier()

# Lists to store metrics for each fold
```

```

RF_precision_list = []
RF_recall_list = []
RF_f1_list = []

for train_index, val_index in stratified_kf.split(X_train, y_train):
    X_train_fold, X_val_fold = X_train.iloc[train_index], X_train.
    ↪iloc[val_index]
    y_train_fold, y_val_fold = y_train.iloc[train_index], y_train.
    ↪iloc[val_index]

    RF.fit(X_train_fold, y_train_fold)
    RF_pred = RF.predict(X_val_fold)

    # Calculate metrics for each fold
    classification_report_fold = classification_report(y_val_fold, RF_pred,
    ↪output_dict=True)

    RF_precision_list.append(classification_report_fold['weighted_
    ↪avg']['precision'])
    RF_recall_list.append(classification_report_fold['weighted avg']['recall'])
    RF_f1_list.append(classification_report_fold['weighted avg']['f1-score'])

# Calculate mean metrics across all folds
mean_RF_precision = np.mean(RF_precision_list)
mean_RF_recall = np.mean(RF_recall_list)
mean_RF_f1 = np.mean(RF_f1_list)

# Print or use the mean metrics as needed
print(f'Mean Precision for RF: {mean_RF_precision}')
print(f'Mean Recall for RF: {mean_RF_recall}')
print(f'Mean F1-Score for RF: {mean_RF_f1}')

```

Mean Precision for RF: 0.743376848791913

Mean Recall for RF: 0.7493508584389676

Mean F1-Score for RF: 0.7307146385847803

KNN

[37]: RF = KNeighborsClassifier()

```

# Lists to store metrics for each fold
RF_precision_list = []
RF_recall_list = []
RF_f1_list = []

for train_index, val_index in stratified_kf.split(X_train, y_train):
    X_train_fold, X_val_fold = X_train.iloc[train_index], X_train.
    ↪iloc[val_index]

```

```

    y_train_fold, y_val_fold = y_train.iloc[train_index], y_train.
    ↪iloc[val_index]

    RF.fit(X_train_fold, y_train_fold)
    RF_pred = RF.predict(X_val_fold)

    # Calculate metrics for each fold
    classification_report_fold = classification_report(y_val_fold, RF_pred,
    ↪output_dict=True)

    RF_precision_list.append(classification_report_fold['weighted_
    ↪avg']['precision'])
    RF_recall_list.append(classification_report_fold['weighted avg']['recall'])
    RF_f1_list.append(classification_report_fold['weighted avg']['f1-score'])

# Calculate mean metrics across all folds
mean_RF_precision = np.mean(RF_precision_list)
mean_RF_recall = np.mean(RF_recall_list)
mean_RF_f1 = np.mean(RF_f1_list)

# Print or use the mean metrics as needed
print(f'Mean Precision for RF: {mean_RF_precision}')
print(f'Mean Recall for RF: {mean_RF_recall}')
print(f'Mean F1-Score for RF: {mean_RF_f1}')

```

Mean Precision for RF: 0.6973526563246721
Mean Recall for RF: 0.7094800146730681
Mean F1-Score for RF: 0.6988606461552431

Naive Bayes

```

[38]: RF = GaussianNB()

# Lists to store metrics for each fold
RF_precision_list = []
RF_recall_list = []
RF_f1_list = []

for train_index, val_index in stratified_kf.split(X_train, y_train):
    X_train_fold, X_val_fold = X_train.iloc[train_index], X_train.
    ↪iloc[val_index]
    y_train_fold, y_val_fold = y_train.iloc[train_index], y_train.
    ↪iloc[val_index]

    RF.fit(X_train_fold, y_train_fold)
    RF_pred = RF.predict(X_val_fold)

```

```

    # Calculate metrics for each fold
    classification_report_fold = classification_report(y_val_fold, RF_pred,
    ↪output_dict=True)

    RF_precision_list.append(classification_report_fold['weighted_
    ↪avg']['precision'])
    RF_recall_list.append(classification_report_fold['weighted avg']['recall'])
    RF_f1_list.append(classification_report_fold['weighted avg']['f1-score'])

# Calculate mean metrics across all folds
mean_RF_precision = np.mean(RF_precision_list)
mean_RF_recall = np.mean(RF_recall_list)
mean_RF_f1 = np.mean(RF_f1_list)

# Print or use the mean metrics as needed
print(f'Mean Precision for RF: {mean_RF_precision}')
print(f'Mean Recall for RF: {mean_RF_recall}')
print(f'Mean F1-Score for RF: {mean_RF_f1}')

```

Mean Precision for RF: 0.6615354006784583

Mean Recall for RF: 0.6781883307616214

Mean F1-Score for RF: 0.6643549622105983

SVM

[39]: RF = SVC()

```

# Lists to store metrics for each fold
RF_precision_list = []
RF_recall_list = []
RF_f1_list = []

for train_index, val_index in stratified_kf.split(X_train, y_train):
    X_train_fold, X_val_fold = X_train.iloc[train_index], X_train.
    ↪iloc[val_index]
    y_train_fold, y_val_fold = y_train.iloc[train_index], y_train.
    ↪iloc[val_index]

    RF.fit(X_train_fold, y_train_fold)
    RF_pred = RF.predict(X_val_fold)

    # Calculate metrics for each fold
    classification_report_fold = classification_report(y_val_fold, RF_pred,
    ↪output_dict=True)

    RF_precision_list.append(classification_report_fold['weighted_
    ↪avg']['precision'])

```

```

RF_recall_list.append(classification_report_fold['weighted avg']['recall'])
RF_f1_list.append(classification_report_fold['weighted avg']['f1-score'])

# Calculate mean metrics across all folds
mean_RF_precision = np.mean(RF_precision_list)
mean_RF_recall = np.mean(RF_recall_list)
mean_RF_f1 = np.mean(RF_f1_list)

# Print or use the mean metrics as needed
print(f'Mean Precision for RF: {mean_RF_precision}')
print(f'Mean Recall for RF: {mean_RF_recall}')
print(f'Mean F1-Score for RF: {mean_RF_f1}')

```

Mean Precision for RF: 0.7479669321680513

Mean Recall for RF: 0.674892186514749

Mean F1-Score for RF: 0.5607237621001302

1.0.11 Plotting the ROC curves for all 6 models

```

[40]: models = {
    'Logistic Regression': LogisticRegression(),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(),
    'KNN': KNeighborsClassifier(),
    'Naive Bayes': GaussianNB(),
    # 'SVM': SVC(probability=True)
}

# Set up k-fold cross-validation
stratified_kf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Iterate through each model
for model_name, model in models.items():
    print(f"Evaluating {model_name}")

    # Perform k-fold cross-validation
    mean_fpr = np.linspace(0, 1, 100)
    tpr_list = []

    for train_index, val_index in stratified_kf.split(X_train, y_train):
        X_train_fold, X_val_fold = X_train.iloc[train_index], X_train.
        ↪iloc[val_index]
        y_train_fold, y_val_fold = y_train.iloc[train_index], y_train.
        ↪iloc[val_index]

        model.fit(X_train_fold, y_train_fold)
        y_pred = model.predict_proba(X_val_fold)[: , 1]

```

```

fpr, tpr, _ = roc_curve(y_val_fold, y_pred)
tpr_list.append(np.interp(mean_fpr, fpr, tpr))

# Compute mean and standard deviation of the ROC curves
mean_tpr = np.mean(tpr_list, axis=0)
mean_auc = auc(mean_fpr, mean_tpr)
std_auc = np.std(tpr_list, axis=0)

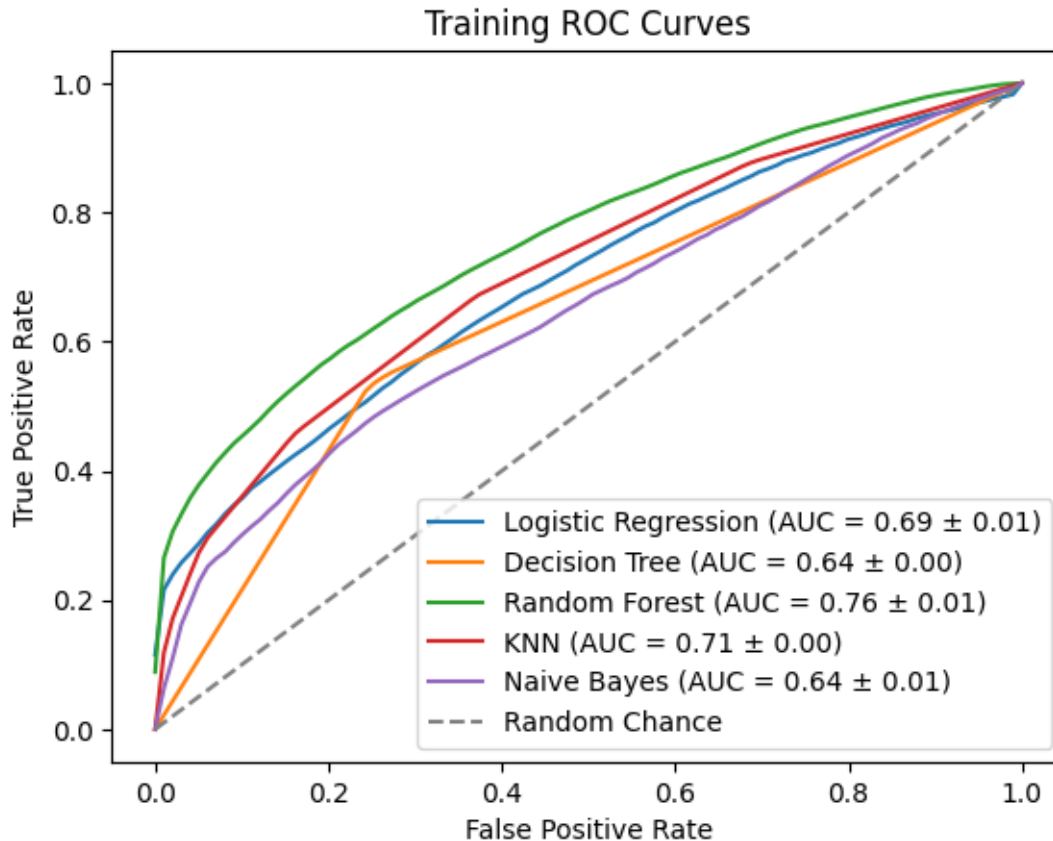
# Plot the training ROC curve with mean and standard deviation
plt.plot(mean_fpr, mean_tpr, label=f'{model_name} (AUC = {mean_auc:.2f} ±_{np.mean(std_auc):.2f})')

# Plot the random chance line
plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random Chance')

# Set plot labels and title
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Training ROC Curves')
plt.legend(loc='lower right')
plt.show()

```

Evaluating Logistic Regression
 Evaluating Decision Tree
 Evaluating Random Forest
 Evaluating KNN
 Evaluating Naive Bayes



1.0.12 select models for ensemble based on roc curves, 3 or all?

Ensemble Method

```
[41]: from sklearn.metrics import log_loss
from sklearn.ensemble import VotingClassifier, StackingClassifier

LR = LogisticRegression()
DT = DecisionTreeClassifier()
RF = RandomForestClassifier()
KNN = KNeighborsClassifier()
SVM = SVC()
NB = GaussianNB()

## Voting Classifier
voting_classifier = VotingClassifier(estimators=[('rf', LR),
                                                ('knn', KNN),
                                                ('lr', LR)],
                                   voting='hard')

# 'hard' for majority voting,
# 'soft' for weighted voting based on probabilities
```

```

# Stacking Classifier
stacking_classifier = StackingClassifier(estimators=[('rf', LR),
                                                    ('knn', KNN),
                                                    ('lr', LR)],
                                       final_estimator=LogisticRegression())

# Train and evaluate each classifier
classifiers = [voting_classifier,
               stacking_classifier]

for classifier in classifiers:
    classifier.fit(X_train, y_train)
    y_pred = classifier.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"{classifier.__class__.__name__} Accuracy: {accuracy}")

```

VotingClassifier Accuracy: 0.7246703991331046

StackingClassifier Accuracy: 0.7338811630847029

1.0.13 Random forest works better, followed by KNN, Logistic Regression, Decision Tree and SVM in that order